Snowflake Data Engineer Exam Questions

Question 1: To view/monitor the clustering metadata for a table, Snowflake provides which of the following system functions?

A) SYSTEM$CLUSTERING_DEPTH_KEY

B) SYSTEM$CLUSTERING_KEY_INFORMATION (including clustering depth)

C) SYSTEM$CLUSTERING_DEPTH

D) SYSTEM$CLUSTERING_INFORMATION (including clustering depth)

Answer: C, D

Explanation: SYSTEM$CLUSTERING_DEPTH: Computes the average depth of the table according to the specified columns...

Question 2: Select the incorrect statements regarding Clustering depth?

A) The clustering depth for a populated table measures the average depth (1 or greater) of the overlapping micro-partitions...

B) It helps Monitoring the clustering 'health' of a large table, particularly over time as DML is performed on the table.

C) Clustering depth can be used for determining whether a large table would benefit from explicitly defining a clustering key.

D) A table with no micro-partitions (i.e. an unpopulated/empty table) has a clustering depth of 1.

Answer: D

Explanation: A table with no micro-partitions (i.e. an unpopulated/empty table) has a clustering depth of 0.

Question 3: While creating even Secure UDF, Snowflake recommends using randomized identifiers instead of sequence-generated values?

A) TRUE

B) FALSE

Answer: A

Explanation: Snowflake recommends using randomized identifiers for better security.

Question 4: Which UDF programming language is not supported with Snowflake Secure Data Sharing feature?

A) SQL

B) JAVA

C) JAVASCRIPT

D) PYTHON

Answer: C

Explanation: JAVASCRIPT is not supported for Snowflake Secure Data Sharing.

Question 5: Which are supported Programming Languages for Creating UDTFs?

A) Python

B) Node.javascript

C) Javascript

D) Java

E) Perl

Answer: A, C, D

Explanation: Python, Javascript, and Java are supported for creating UDTFs.

Question 6: Tasks may optionally use table streams to provide a convenient way to continuously process new or changed data. Which System Function can be used by Data Engineer to verify whether a stream contains changed data for a table?

A) SYSTEM$STREAM_HAS_CHANGE_DATA

B) SYSTEM$STREAM_CDC_DATA

C) SYSTEM$STREAM_HAS_DATA

D) SYSTEM$STREAM_DELTA_DATA

Answer: C

Explanation: SYSTEM$STREAM_HAS_DATA indicates whether a specified stream contains change data capture records.

Question 7: Streams cannot be created to query change data on which of the following objects? [Select All that Apply]

A) Standard tables, including shared tables.

B) Views, including secure views

C) Directory tables

D) Query Log Tables

E) External tables

Answer: D

Explanation: Streams do not support Query Log tables.

Question 8: To advance the offset of a stream to the current table version without consuming the change data in a DML operation, which of the following operations can be done by Data Engineer?

A) Using the CREATE OR REPLACE STREAM syntax, recreate the STREAM

B) Insert the current change data into a temporary table...

C) A stream advances the offset only when it is used in a DML transaction...

D) Delete the offset using STREAM properties SYSTEM$RESET_OFFSET(<stream_id>)

Answer: A, B

Explanation: Recreating the stream or inserting the current change data into a temporary table will advance the offset.